

Applied Data Science and Machine Learning Course

Mohammad Sabik Irbaz

Lead Machine Learning Engineer
Pioneer Alpha Ltd.
sabik@pioneeralpha.com



Powered by Pioneer Alpha Ltd.

TABLE OF CONTENTS

01	Introduction
02	What is Machine Learning (ML) ?
03	Machine Learning: The Origins
04	Machine Learning: State of the Art
05	Do we need traditional ML in the deep learning era?
06	ML Production and Data Science
07	Types of ML Algorithms
08	Cost Function, Gradient Descent & Linear Regression
09	Concluding Remarks

Introduction




Pioneer Alpha Ltd. is an emerging R&D based Computer Software Company which provides IT-enabled products and services rendered by internet or electronic network.

Introduction



Pioneer Alpha Ltd. is an emerging R&D based Computer Software Company which provides IT-enabled products and services rendered by internet or electronic network.



Amar iSchool is an online interactive educational platform for learners. You can increase your expertise level with Amar iSchool by watching courses and interacting with the peers and experts anytime, anywhere.

“Applied Data Science and Machine Learning”

a course by Amar iSchool

02

Applied Data Science and Machine Learning

From the amount of data that is being produced every day, its use in business, politics, administration, or social work is increasing day by day. The arrival of 5G will further speed up the amount of this data. Then we may have to work with big data all the time. Currently, the global job growth rate in data science is around 650%. Python is like the Swiss Army Knife of technology; Python solves any task starting from the web. So if you don't enroll in this fantastic combo course in Python and Data Science, you will regret later.

BEGINNER ★★★★★ 0 (0 Ratings) 10 Students enrolled

Created by Mridul Banik Last updated Tue, 09-Feb-2021 Bengali

“Applied Data Science and Machine Learning”

a course by Amar iSchool

02

Applied Data Science and Machine Learning

From the amount of data that is being produced every day, its use in business, politics, administration, or social work is increasing day by day. The arrival of 5G will further speed up the amount of this data. Then we may have to work with big data all the time. Currently, the global job growth rate in data science is around 650%. Python is like the Swiss Army Knife of technology; Python solves any task starting from the web. So if you don't enroll in this fantastic combo course in Python and Data Science, you will regret later.

BEGINNER



(0 Ratings) 10 Students enrolled

Created by Mridul Banik Last updated Tue, 09-Feb-2021 Bengali

What will i learn?

- ✓ Learn Advanced Python3
- ✓ Data mining and wrangling
- ✓ How to use power BL for Data Exploration
- ✓ Data Modeling and Validation for supervised and unsupervised Machine Learning
- ✓ Deploying Predictive Models to Production
- ✓ Data scrapping (From web page / JSON based API)
- ✓ Exploratory Data analysis
- ✓ Data Pre-processing
- ✓ Data Science in Cloud

“Applied Data Science and Machine Learning”

a course by Amar iSchool

02

Applied Data Science and Machine Learning

From the amount of data that is being produced every day, its use in business, politics, administration, or social work is increasing day by day. The arrival of 5G will further speed up the amount of this data. Then we may have to work with big data all the time. Currently, the global job growth rate in data science is around 650%. Python is like the Swiss Army Knife of technology; Python solves any task starting from the web. So if you don't enroll in this fantastic combo course in Python and Data Science, you will regret later.

BEGINNER ★★★★★ 0 (0 Ratings) 10 Students enrolled

Created by Mridul Banik Last updated Tue, 09-Feb-2021 Bengali

What will i learn?

- ✓ Learn Advanced Python3
- ✓ Data mining and wrangling
- ✓ How to use power BI for Data Exploration
- ✓ Data Modeling and Validation for supervised and unsupervised Machine Learning
- ✓ Deploying Predictive Models to Production
- ✓ Data scrapping (From web page / JSON based API)
- ✓ Exploratory Data analysis
- ✓ Data Pre-processing
- ✓ Data Science in Cloud

Curriculum for this course		3 Lessons	00:00:00 Hours
-	Course Inauguration		00:00:00 Hours
+	Environment Setup	2 Lessons	00:00:00 Hours
+	Python Programming for Data Science	1 Lessons	00:00:00 Hours
+	Diving into Data Science	0 Lessons	00:00:00 Hours
+	Data Collection	0 Lessons	00:00:00 Hours
+	Diving into Data Exploration	0 Lessons	00:00:00 Hours
+	Diving into Data Preprocessing	0 Lessons	00:00:00 Hours
+	Power BI, Tableau, Introduction of KNIME	0 Lessons	00:00:00 Hours
+	Diving into Data Modeling and Validation	0 Lessons	00:00:00 Hours
+	Data Science in Cloud	0 Lessons	00:00:00 Hours
+	Deploying Predictive Models to Production	0 Lessons	00:00:00 Hours

“Applied Data Science and Machine Learning”

a course by Amar iSchool

02

Applied Data Science and Machine Learning

From the amount of data that is being produced every day, its use in business, politics, administration, or social work is increasing day by day. The arrival of 5G will further speed up the amount of this data. Then we may have to work with big data all the time. Currently, the global job growth rate in data science is around 650%. Python is like the Swiss Army Knife of technology; Python solves any task starting from the web. So if you don't enroll in this fantastic combo course in Python and Data Science, you will regret later.

BEGINNER ★★★★★ 0 (0 Ratings) 10 Students enrolled

Created by Mridul Banik Last updated Tue, 09-Feb-2021 Bengali

What will i learn?

- ✓ Learn Advanced Python3
- ✓ Data mining and wrangling
- ✓ How to use power BI for Data Exploration
- ✓ Data Modeling and Validation for supervised and unsupervised Machine Learning
- ✓ Deploying Predictive Models to Production
- ✓ Data scrapping (From web page / JSON based API)
- ✓ Exploratory Data analysis
- ✓ Data Pre-processing
- ✓ Data Science in Cloud

Curriculum for this course		3 Lessons	00:00:00 Hours
-	Course Inauguration		00:00:00 Hours
+	Environment Setup	2 Lessons	00:00:00 Hours
+	Python Programming for Data Science	1 Lessons	00:00:00 Hours
+	Diving into Data Science	0 Lessons	00:00:00 Hours
+	Data Collection	0 Lessons	00:00:00 Hours
+	Diving into Data Exploration	0 Lessons	00:00:00 Hours
+	Diving into Data Preprocessing	0 Lessons	00:00:00 Hours
+	Power BI, Tableau, Introduction of KNIME	0 Lessons	00:00:00 Hours
+	Diving into Data Modeling and Validation	0 Lessons	00:00:00 Hours
+	Data Science in Cloud	0 Lessons	00:00:00 Hours
+	Deploying Predictive Models to Production	0 Lessons	00:00:00 Hours

Requirements:

- Basics about Coding
- Computer Fundamentals

“Applied Data Science and Machine Learning”

a course by Amar iSchool (cont.)

02

More stats:

- 50+ hours recorded classes
- 20+ hours live classes
- Documentation
- 15+ ML Algorithm based projects
- Real life project
- 15+ assignments and solve in live
- Project Deployment in Cloud
- Project Deployment using Flask
- Project API Creation
- Hyperparameter tuning on Existing GitHub Work
- Debugging
- Mock Test for Job Interview
- Personal Code Review
- Good performance lead to internship opportunities in Pioneer Alpha

TABLE OF CONTENTS

01 Introduction

02 What is Machine Learning (ML) ?

03 Machine Learning: The Origins

04 Machine Learning: State of the Art

05 Do we need traditional ML in the deep learning era?

06 ML Production and Data Science

07 Types of ML Algorithms

08 Cost Function, Gradient Descent & Linear Regression

09 Concluding Remarks

Machine Learning

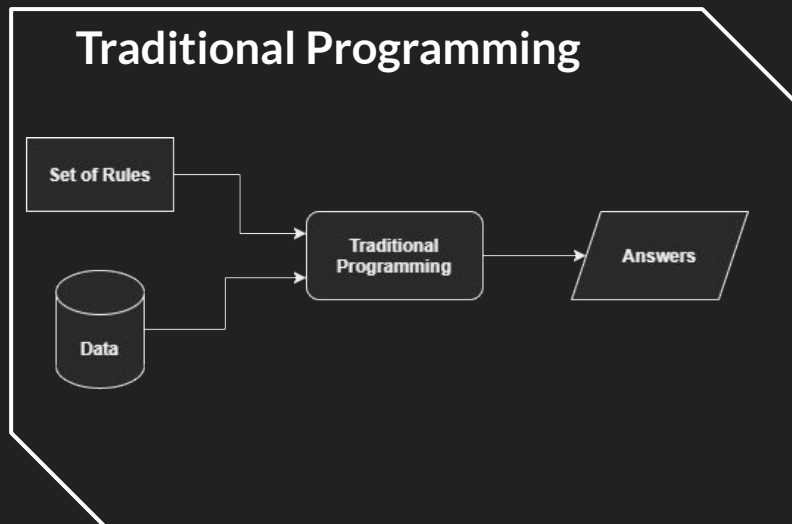
According to Arthur Samuel, one of the pioneers of AI

“Machine Learning is the field of study that gives the computers the ability to learn without being explicitly programmed.” [1959]

Machine Learning

According to Arthur Samuel, one of the pioneers of AI

“Machine Learning is the field of study that gives the computers the ability to learn without being explicitly programmed.” [1959]




Machine Learning


According to Arthur Samuel, one of the pioneers of AI

“Machine Learning is the field of study that gives the computers the ability to learn without being explicitly programmed.” [1959]


Traditional Programming



Machine Learning




AI vs ML vs DL




AI vs ML vs DL

Programs with the
ability to learn and
reason like humans



AI vs ML vs DL


Programs with the ability to learn and reason like humans



Algorithms with the ability to learn without being explicitly programmed

AI vs ML vs DL

Programs with the ability to learn and reason like humans



Algorithms with the ability to learn without being explicitly programmed

Subset of ML in which Artificial Neural Networks adapt and learn from vast amounts of data

TABLE OF CONTENTS

- 01 Introduction**
- 02 What is Machine Learning (ML) ?**
- 03 Machine Learning: The Origins**
- 04 Machine Learning: State of the Art**
- 05 Do we need traditional ML in the deep learning era?**
- 06 ML Production and Data Science**
- 07 Types of ML Algorithms**
- 08 Cost Function, Gradient Descent & Linear Regression**
- 09 Concluding Remarks**

The Origin

Arthur Samuel from IBM developed a computer program for playing checkers in the 1950s which started to beat good checkers players.

He first came up with the phrase “Machine Learning” in 1952.



Taken from Google Arts and Culture

The Origin (cont.)

Perceptron [1957]

Frank Rosenblatt combined Hebb's brain cell model with Samuel's ML efforts and created the perceptron.

Nearest Neighbor Algorithm [1967]

The nearest neighbor algorithm was conceived, which was the beginning of basic pattern recognition.


Multilayer Perceptron

There are many opinions about this. But multilayer perceptron with backpropagation was first popularized by Geoffrey Hinton.

TABLE OF CONTENTS

- 01 Introduction**
- 02 What is Machine Learning (ML) ?**
- 03 Machine Learning: The Origins**
- 04 Machine Learning: State of the Art**
- 05 Do we need traditional ML in the deep learning era?**
- 06 ML Production and Data Science**
- 07 Types of ML Algorithms**
- 08 Cost Function, Gradient Descent & Linear Regression**
- 09 Concluding Remarks**

Four Horsemen of Deep Learning



Geoffrey Hinton

invented backpropagation
won Turing Award 2018


Professor, University of Toronto
Computer scientist, Google Brain



Yann Lecun

invented convolution nets
won Turing Award 2018

VP and Chief AI Scientist,
Facebook



Yoshua Bengio

highest citation in the world in 2018
won Turing Award 2018

Co-Chair of the AI Advisory Council,
government of Canada



Andrew Ng

"democratize deep learning"
"DL with GPU revolution"

Professor, Stanford University
Founder, Google Brain

Quick Question

Can Yann LeCun sing a song?



<https://twitter.com/ylecun/status/1371007395577667585>

Transfer Learning

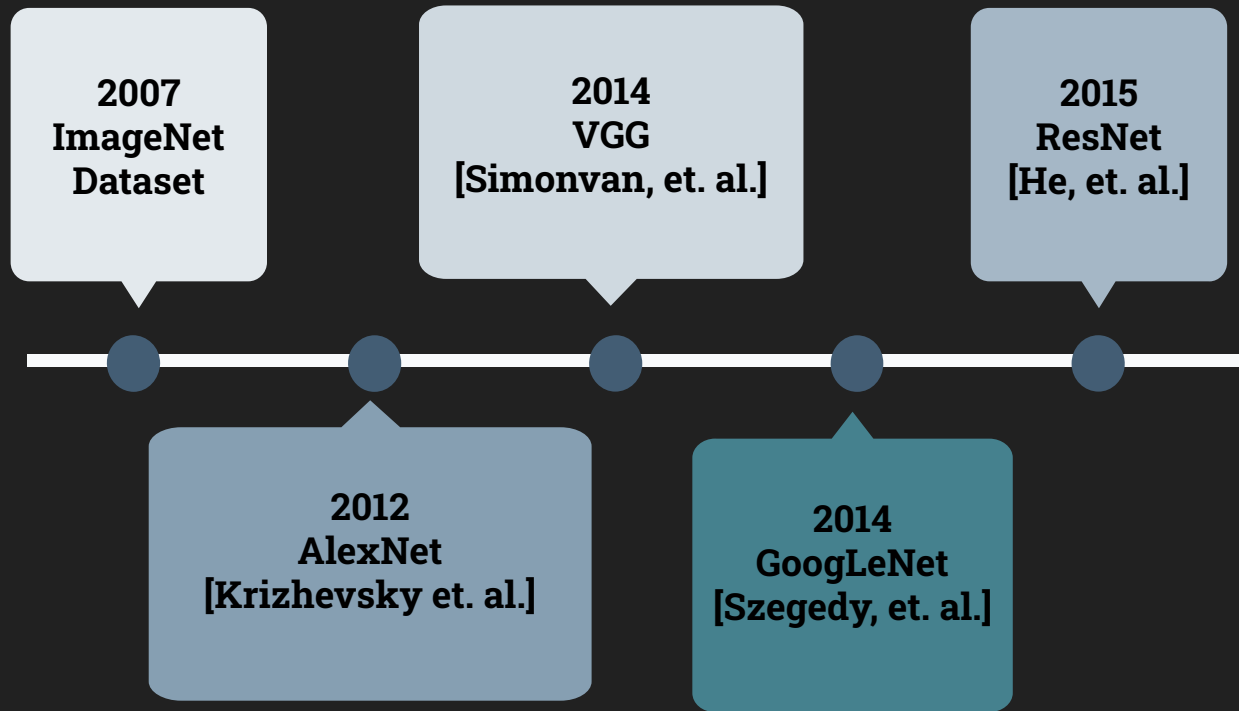
Fine-tuning pre-trained neural networks using custom layers on top for specific tasks.

Transfer Learning


Fine-tuning pre-trained neural networks using custom layers on top for specific tasks.



Computer Vision Revolution [2012 – 2016]



Natural Language Processing (NLP) Revolution [2017 – ongoing]



The Age of Transformers [ongoing]

Vision Transformer (ViT)
[Dosovitskiy et. al. 2020]

GANsformer
[Hudson et. al. 2021]

**Image Segmentation
With Transformer**
[Karimi et. al. 2021]

DeTr for Object Detection
[Carion et. al. 2020]

The Age of Transformers [ongoing]

Vision Transformer (ViT)
[Dosovitskiy et. al. 2020]

GANsformer
[Hudson et. al. 2021]


Image Segmentation
With Transformer
[Karimi et. al. 2021]

DeTr for Object Detection
[Carion et. al. 2020]

DALL-E [OpenAI 2021]

an illustration of a baby daikon radish in a tutu walking a dog

AI-GENERATED IMAGES



**Let's have a look at what Tom
Cruise wants to say us.**

<https://twitter.com/mckaywrigley/status/1364972080513515520>

**Let's have a look at what Tom
Cruise wants to say us.**

<https://twitter.com/mckaywrigley/status/1364972080513515520>

Is AI getting dangerous?

<https://ai.facebook.com/blog/deepfake-detection-challenge-results-an-open-initiative-to-advance-ai/>



TABLE OF CONTENTS

- 01 Introduction
- 02 What is Machine Learning (ML) ?
- 03 Machine Learning: The Origins
- 04 Machine Learning: State of the Art
- 05 Do we need traditional ML in the deep learning era?
- 06 ML Production and Data Science
- 07 Types of ML Algorithms
- 08 Cost Function, Gradient Descent & Linear Regression
- 09 Concluding Remarks

What do the experts say?



Abhisek Thakur
4x Kaggle Grandmaster

`"We should always try baseline models first."`

`https://twitter.com/abhi1thakur/status/1369719373464748040`

Jeremy Howard
Founder of fast.ai

`"Baseline models work better than
your fancy models most of the time."`



`https://twitter.com/omarsar0/status/1370542107232964608`

This is still happening




TABLE OF CONTENTS

- 01 Introduction
- 02 What is Machine Learning (ML) ?
- 03 Machine Learning: The Origins
- 04 Machine Learning: State of the Art
- 05 Do we need traditional ML in the deep learning era?
- 06 ML Production and Data Science
- 07 Types of ML Algorithms
- 08 Cost Function, Gradient Descent & Linear Regression
- 09 Concluding Remarks

ML Production

Chip Huyen (Teaches MLSys at Stanford University)

ML in production: expectation

1. Collect data
2. Train model
3. Deploy model
- 4.



ML Production

Chip Huyen (Teaches MLSys at Stanford University)

ML in production: expectation

1. Collect data
2. Train model
3. Deploy model
- 4.



ML in production: reality

1. Choose a metric to optimize
2. Collect data
3. Train model
4. Realize many labels are wrong -> relabel data
5. Train model
6. Model performs poorly on one class -> collect more data for that class
7. Train model
8. Model performs poorly on most recent data -> collect more recent data
9. Train model
10. Deploy model
11. Dream about \$\$\$
12. Wake up at 2am to complaints that model biases against one group
-> revert to older version
13. Get more data, train more, do more testing
14. Deploy model
15. Pray
16. Model performs well but revenue not increases -> choose a different metric
17. Cry
18. Start over

Data Science



Statistics



Data Science

TABLE OF CONTENTS

- 01 Introduction
- 02 What is Machine Learning (ML) ?
- 03 Machine Learning: The Origins
- 04 Machine Learning: State of the Art
- 05 Do we need traditional ML in the deep learning era?
- 06 ML Production and Data Science
- 07 Types of ML Algorithms
- 08 Cost Function, Gradient Descent & Linear Regression
- 09 Concluding Remarks

Types of ML Algorithms

- Supervised Learning
- Unsupervised Learning
- Semi-Supervised Learning
- Self-Supervised Learning (SSL)
- Reinforcement Learning (RL)
- Active Learning

Supervised vs Unsupervised Learning


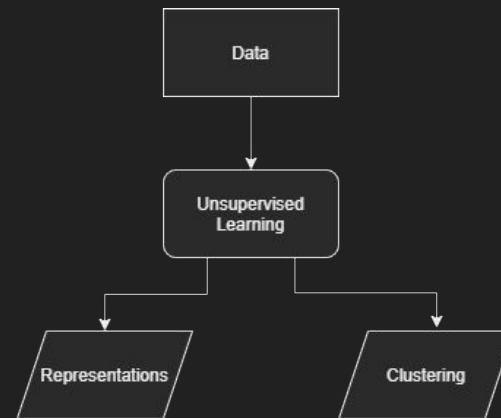

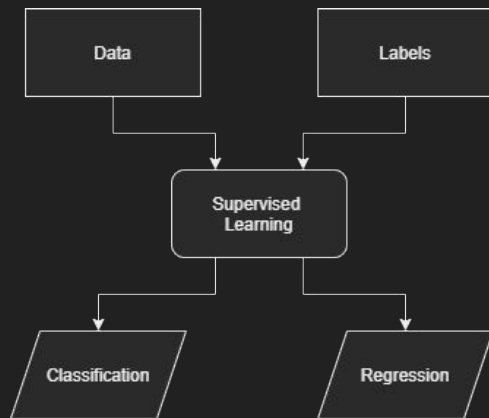


TABLE OF CONTENTS

- 01 Introduction
- 02 What is Machine Learning (ML) ?
- 03 Machine Learning: The Origins
- 04 Machine Learning: State of the Art
- 05 Do we need traditional ML in the deep learning era?
- 06 ML Production and Data Science
- 07 Types of ML Algorithms
- 08 Cost Function, Gradient Descent & Linear Regression
- 09 Concluding Remarks

Hypothesis


Representation of a line : $y = mx + c$

Hypothesis


Representation of a line : $y = mx + c$

Let's redefine it

$$h_{\theta}(x) = \theta_0 + \theta_1 * x$$



Mean Squared Error (MSE)



=> Imagine you have 'm' data points

=> You drew a line

=> MSE calculates how far is that line from data points in average.

=> Mean is for the average

=> Square is for the Euclidean Distance

Cost Function

Cost Function in case of Linear Regression is nothing but MSE

$$J(\theta_0, \theta_1) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

m = #samples

$$h_{\theta}(x) = \theta_0 + \theta_1 * x$$

Cost Function

Cost Function in case of Linear Regression is nothing but MSE

$$J(\theta_0, \theta_1) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

m = #samples

$$h_{\theta}(x) = \theta_0 + \theta_1 * x$$

Our Objective is to minimize the cost.

How do we do that?

Gradient Descent



Gradient Descent (cont.)

Visualization by
Lili Jiang, Head of Data Science, Quora

Gradient Descent (cont.)

Repeat until Convergence

$$\theta_j = \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

Gradient Descent (cont.)

Repeat until Convergence

$$\theta_j = \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

Learning Rate

Gradient Descent (cont.)

Repeat until Convergence

$$\theta_j = \theta_j - \alpha * \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

Learning Rate

Partial Derivative of
the Cost Function

Code Walkthrough

TABLE OF CONTENTS

- 01 Introduction**
- 02 What is Machine Learning (ML) ?**
- 03 Machine Learning: The Origins**
- 04 Machine Learning: State of the Art**
- 05 Do we need traditional ML in the deep learning era?**
- 06 ML Production and Data Science**
- 07 Types of ML Algorithms**
- 08 Cost Function, Gradient Descent & Linear Regression**
- 09 Concluding Remarks**

References

1. Kiaming He, et. al. "Deep residual learning for image recognition". in CVPR 2016.
2. Jacob Devlin, Ming-Wei Chang. "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding". in NAACL-HLT 2018.
3. Ashish Vaswani, Noam Shazeer. "Attention is all you need". in NeurIPS 2017
4. Chip Huyen. "Machine Learning Systems Design"
5. Andrew Ng. "Machine Learning"
6. Fei-Fei Li. "Convolutional Neural Networks for Visual Recognition"
7. Chris Manning. "Natural Language Processing with Deep Learning"